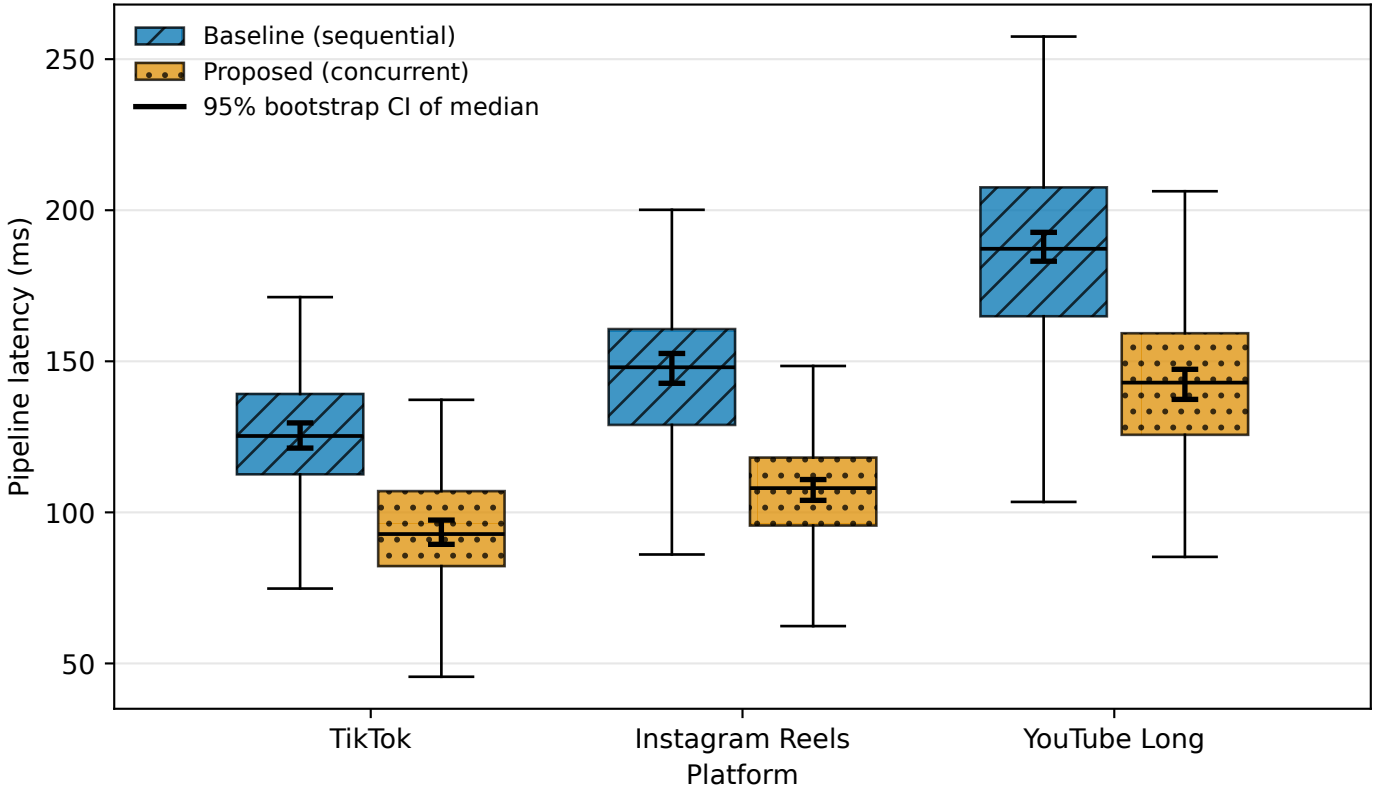
Per-run latency distribution (evaluation phase, warm-up filtered)
Baseline (sequential)
Proposed (concurrent)
95% bootstrap CI of median
Pipeline latency (ms)
250
200
150
100
50
TikTok
Instagram Reels
YouTube Long
Platform